


Chapter 4: Threads & Concurrency






Outline


- Overview
- Multicore Programming
- Multithreading Models
- Thread Libraries
- Implicit Threading
- Threading Issues
- Operating System Examples






Objectives


- Identify the basic components of a thread, and contrast threads and processes
- Describe the benefits and challenges of designing multithreaded applications
- Illustrate different approaches to implicit threading including thread pools, fork-join, and Grand Central Dispatch
- Describe how the Windows and Linux operating systems represent threads
- Designing multithreaded applications using the Pthreads, Java, and Windows threading APIs





Motivation

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels are generally multithreaded



Aplicatii concurente

- proces = un singur punct de executie in aplicatie (o singura instanta in rulare a aplicatiei)
- unele aplicatii pot profita de existenta mai multor puncte de executie simultana in cadrul aplicatiei (mai ales pe multiprocesoare)
- ex. 1: procesor cu 4 core-uri + aplicatie de filtrare de imagini care imparte imaginea in 4 cadrane, fiecare core filtrand un cadran
 - avantaj: descompunerea activitatilor mari in activitati mai simple care ruleaza simultan => reducerea timpului de rulare
- ex. 2: un singur procesor + program de gestiune a ferestrelor in GUI
 - verifica si proceseaza input-ul de la tastatura, mouse si retea pt. a produce output pe ecran
 - conceptual, exista mai multe activitati concurente: verificare/procesare input tastatura, verificare/procesare input mouse, verificare/procesare input retea, afisare bitmap-uri pe ecran
 - toate aceste activitati sunt codate in module separate, cu date private si comunica intre ele prin data partajate

Aplicatii concurente (cont'd)

- ex. 3: server de retea
 - activitatea principala: asteapta cereri de la client, le proceseaza si trimite inapoi raspunsurile
 - cu un singur punct de executie procesarea unei cereri particulare ia mult timp (de pilda, in asteptarea datelor de pe disc) => se intarzie foarte mult tratarea altor cereri (situatie similara cu cea care a condus la nevoia de multitasking)
 - solutie: fiecare cerere client e procesata in alt punct de executie al aplicatiei server

Procese si date partajate

- crearea de puncte multiple de executie in program se poate face cu procese care partajeaza date (sa zicem prin memorie partajata, cu acces protejat cu semafoare/locks)
 - procesele au date private (datorita protectiei MMU a spatiului de adresa)
 - comunica intre ele prin IPC (memorie partajata)
 - daca exista capacitate de multiprocesare, procesele pot rula simultan pe mai multe procesoare
- puncte nevralgice
 - crearea proceselor
 - context-switch-ul

Crearea proceselor


- contextul unui proces e stufos (stocat in PCB in kernel)
- include starea complete a CPU, registre de gestiunea memoriei, tabele de pagini de memorie, descriptori de fisiere, actiuni asociate semnalelor, etc

=> cost semnificativ de creare a unui proces

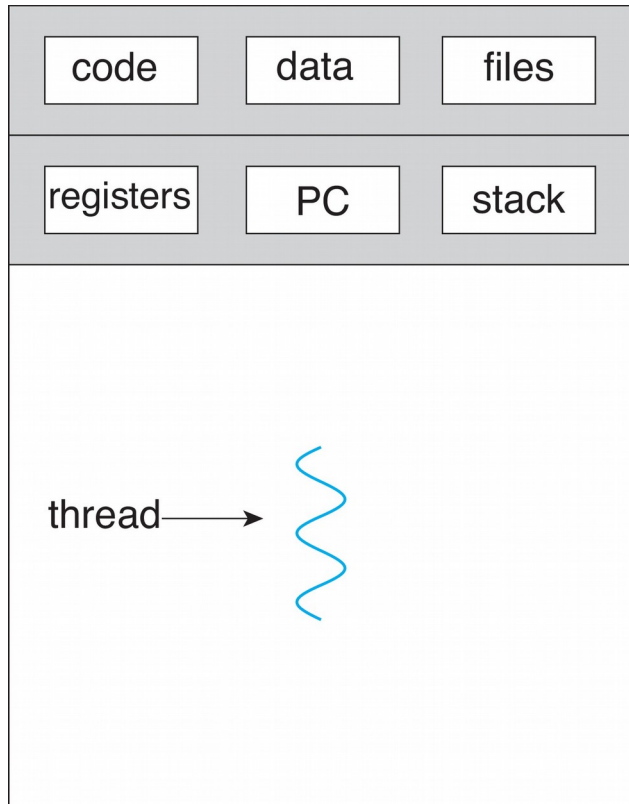
- daca acest cost e prea mare, aplicatiile pot sa nu foloseasca eficient procesoarele:
 - ex aplicatie de filtrare imagini
 - daca timpul de creare a unui proces e comparabil cu timpul de filtrare al unui cadran, e posibil ca filtrarea intregii imagini cu 4 procese sa dureze mai mult decat filtrarea ei cu un singur proces ("slowdown")

Context-switch


- salvarea contextul unui proces si incarcarea contextului unui nou proces poate implica un cost semnificativ avand in vedere bogatia de informatii din PCB
- ex: secventa de evenimente pt. producator-consumator cu zero buffering (“rendez-vous”)
 - producatorul produce un element si se blocheaza
 - sistemul face context-switch si aduce consumatorul pe procesor
 - consumatorul consuma elementul si se blocheaza
 - sistemul face context-switch si aduce producatorul inapoi pe procesor pt a produce un nou element
- daca timpul de context-switch e mare, viteza de transfer a datelor intre producator si consumator e serios afectata (fiecare transfer implica doua context-switch-uri); in orice caz, e imposibil sa se atinga viteza teoretica maxima de transfer



Single and Multithreaded Processes



single-threaded process

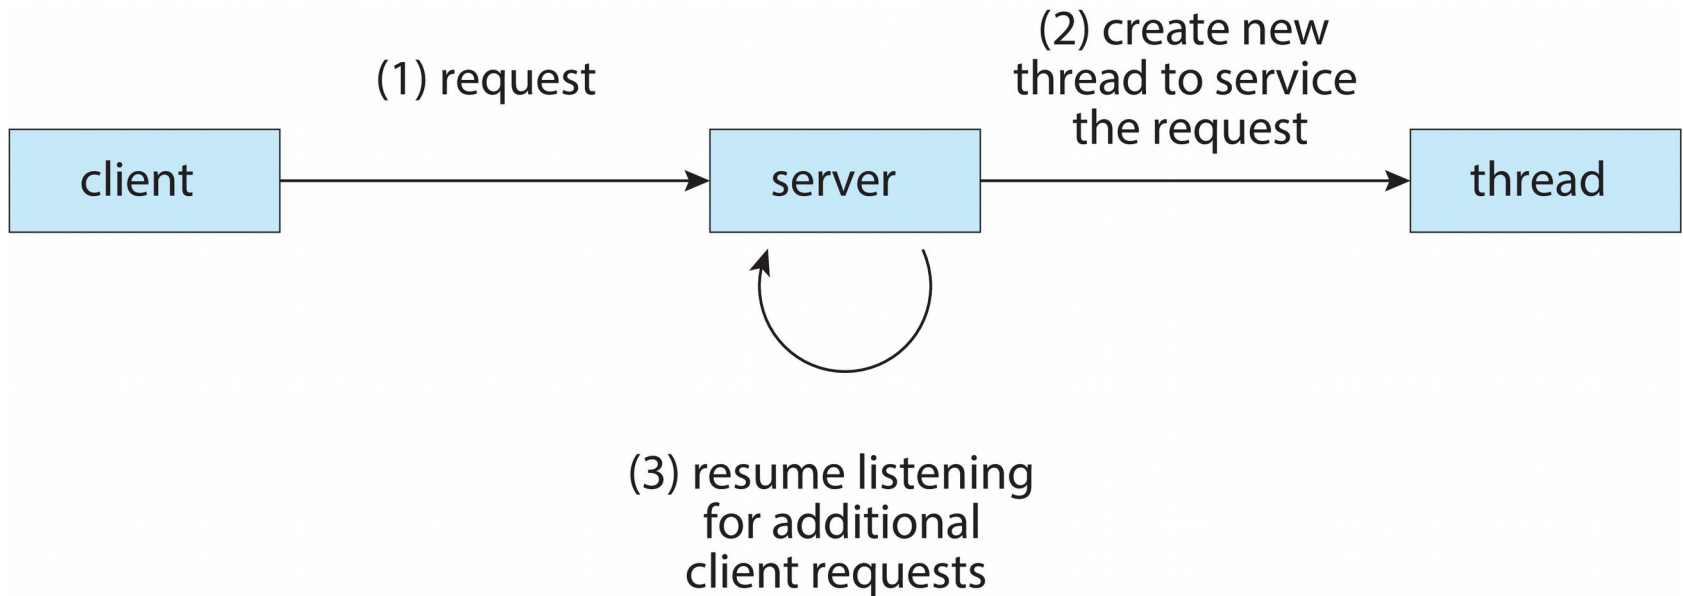



multithreaded process






Multithreaded Server Architecture





Benefits


- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- **Scalability** – process can take advantage of multicore architectures





Multicore Programming


- **Multicore** or **multiprocessor** systems puts pressure on programmers, challenges include:
 - **Dividing activities**
 - **Balance**
 - **Data splitting**
 - **Data dependency**
 - **Testing and debugging**
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
 - Single processor / core, scheduler providing concurrency






Concurrency vs. Parallelism

- **Concurrent execution on single-core system:**




- **Parallelism on a multi-core system:**






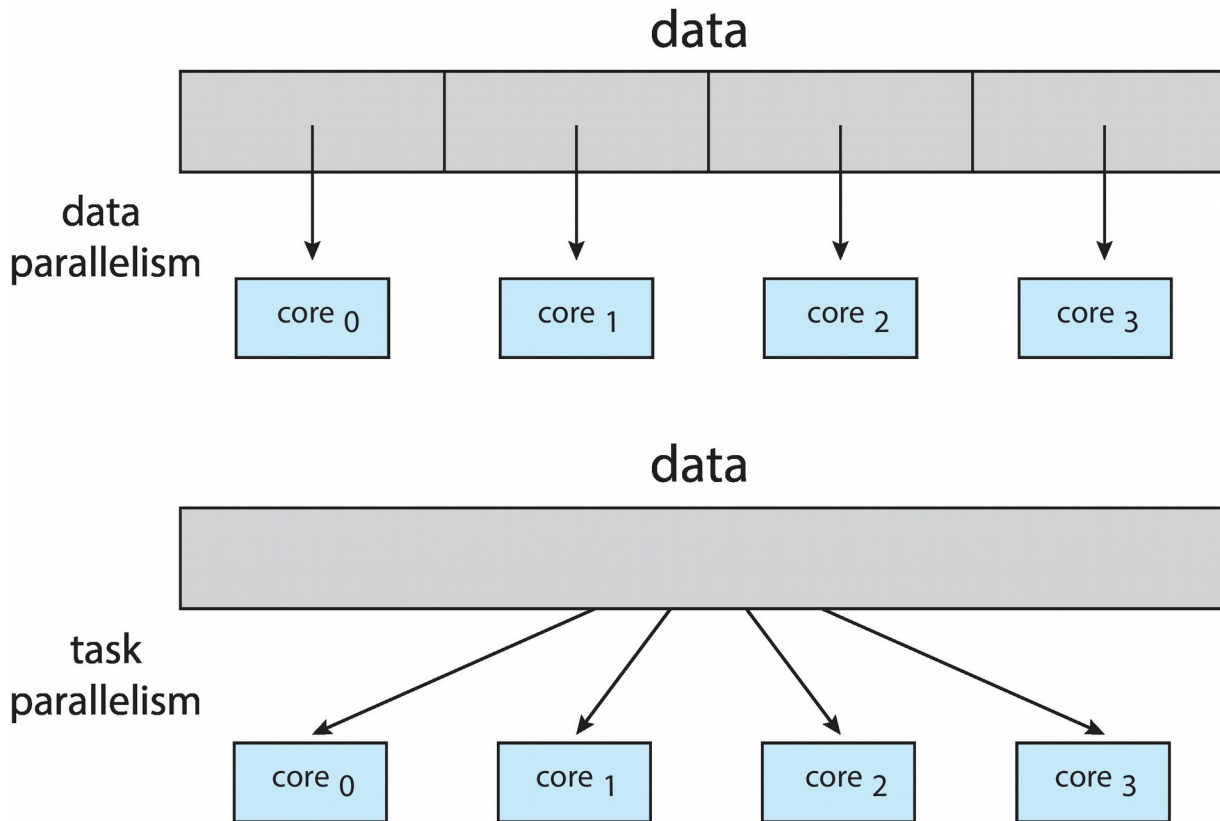
Multicore Programming


- Types of parallelism
 - **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
 - **Task parallelism** – distributing threads across cores, each thread performing unique operation





Data and Task Parallelism





Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components
- S is serial portion
- N processing cores


$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches $1 / S$


Serial portion of an application has disproportionate effect on performance gained by adding additional cores

- But does the law take into account contemporary multicore systems?





Amdahl's Law



Legea lui Gustafson

- Amdahl

$$Speedup = 1 / (s + p / N)$$

- presupunere implicita: p independent de N (i.e., dimensiunea problemei e fixa)
- presupunere nerealista in practica unde dimensiunea problemei scaleaza cu numarul de procesoare => o noua abordare, in care se pp ca timpul de rulare e constant, nu dimensiunea problemei
- acum, se considera ca p variaza liniar cu N
- observatie practica: partea seriala s (initializare vectori, incarcare program, bottleneck-uri seriale, I/O) nu creste odata cu dimensiunea problemei




FIGURE 2a. Fixed-Size Model: $Speedup = 1 / (s + p / N)$

Legea lui Gustafson (cont'd)

- definitie noua pt speedup

$$\text{Scaled speedup} = s + pN = s + (1 - s)N = N + (1 - N)s$$

- daca s e constant, dependenta de N e liniara cu panta $1 - s$

=> schimbare de paradigma: scaled speedup-ul este de fapt “slowdown-ul” (incetinirea) teoretica pe care o inregistreaza un program paralel atunci cand ar rula ipotetic pe o masina cu un singur procesor




FIGURE 2b. Scaled-Size Model: $\text{Speedup} = s + Np$

Concluzii legea lui Gustafson

- accentul folosirii unui nr crescut de procesoare cade pe necesitatea de a rula programe mai mari in acelasi timp daca se poate, nu pe scaderea timpului de executie al unei probleme de dimensiune fixa
- obs: nu orice problema poate fi marita arbitrar, ca atare exista si limitari ale legii lui Gustafson

Fire de executie (threads of execution)

- modalitate de a reduce costurile crearii punctelor de executie in aplicatie si a schimbarii contextelor de executie intre ele
- idee: puncte de executie multiple din aplicatie partajeaza o parte din contextul programului (registrele de gestiune a memoriei, tabelele de pagini, descriptorii de fisiere deschise, samd)
- DAR, fiecare punct de executie are:
 - o copie individuala a unui subset al contextului de executie a aplicatiei (de ex. contextul CPU)
 - structuri de date necesare punctului de executie (de ex. stiva)
- apare o noua abstractie de nivel inalt, *firul de executie (thread)* = un punct de executie cu context redus in cadrul programului
 - referit uneori din cauza acestei viziuni ca fiind un “proces usor” (lightweight process, LWP)

Caracteristici threaduri

- ruleaza secvential, au program counter si stiva propria
- multiplexeaza accesul la CPU ca si procesele (pe multiprocesoare ruleaza in paralel)
- pot crea alte threaduri
- pot executa apeluri de sistem
 - daca un thread se blocheaza intr-un apel sistem, alt thread primeste procesorul
- analogie posibila
 - threadul este fata de proces ceea ce procesul e in raport cu procesorul
 - procesul actioneaza ca un procesor virtual pe care ruleaza thread-ul

Diferente fata de procese

- threadurile aceluiasi proces partajeaza spatiul de adresa al procesului (de ex, partajeaza variabilele globale) => un thread poate distruge usor alt thread (nu exista protectia MMU ca in cazul proceselor diferite)
- lipsa protectiei intre threaduri
 - e inevitabila prin design
 - nu e necesara (threadurile sunt parte a aceluiasi program al unui anumit utilizator)
 - nici nu e de dorit (impunerea unor domenii de protectie conduce la probleme similare proceselor)
- alte resurse partajate: acelasi set de fisiere deschise, timere, semnale, etc

Alte caracteristici ale threadurilor

- stari (la fel ca la procese): in rulare, gata de rulare, blocat, terminat
- modele de utilizare
 - cooperativ, lucru in echipa (ex filtrare de imagini)
 - master-slave/worker (ex server)
 - pipeline (ex producator-consumator)
- avantaj principal: datele partajate sunt datele globale din proces (nu e nevoie de setarea unor mecanisme IPC de tipul memoriei partajate)
 - buffer global pt. producator-consumator
 - argument puternic pt sistemele multiprocesor unde threadurile pot rula pe CPU-uri diferite => partajarea implicita a datelor prin spatiul comun de adresa (nu e nevoie de mecanisme speciale de partajare ca in cazul proceselor, eg memorie partajata IPC)

Design-ul pachetelor de threaduri

- pachet de threaduri = colectie de primitive (apeluri de biblioteca) pt lucrul cu thread-uri
- (1) gestiunea threadurilor
 - creare thread: primeste ca argumente functia care reprezinta punctul de executie initial, o stiva privata si o prioritate de planificare pe procesor; intoarce un TID (Thread ID)
 - terminare thread: explicit prin apel "exit" sau semnal de tip kill de la alt thread/proces
 - primitive pt. mecanisme de sincronizare, necesare datorita existentei datelor partajate (uzual mutex-uri, dar in mod notabil si variabile de conditie folosite in conexiune cu un mutex)
 - ex: acquire/release resource folosind mutex+condition variable
- (2) planificare (aceeasi algoritmi ca si la procese, vom discuta la planificarea proceselor/threadurilor)

Design-ul pachetelor de threaduri

- (3) probleme de reentranta
 - scenariu: doua threaduri T1 si T2 executa concurent apeluri sistem, T1 reuseste, T2 esueaza
 - daca T1 nu evalueaza *errno* inainte ca T2 sa execute apelul sistem care esueaza, T1 va crede eronat ca apelul sau sistem a esuat
 - problema principiala: *errno* e variabila globala
 - solutii posibile: (a) protejarea *errno* cu mutex-uri
 - (b) crearea unei copii private a lui *errno* prin intermediul unor variabile globale thread-ului apelant, dar private (invizibile) celorlalte thread-uri => TLS (Thread-Local Storage)

Implementarea threadurilor kernel

- pachetele de threaduri se pot implementa in kernel sau in spatiul utilizator
- threadurile kernel separa campurile din PCB care ajuta la crearea unui punct de executie si le stocheaza intr-un TCB
- astfel, un proces cu un singur punct de executie e reprezentat in kernel de un PCB si un TCB
- operatia de creare a unui thread (*thread_create*) este un apel sistem
 - aloca un TCB
 - aloca stive kernel si user
 - le leaga la PCB-ul procesului in care s-a facut apelul sistem

Exemplu cu doua threaduri kernel

- un thread suspendat in kernel
- altul ruland in spatiul utilizator
- pt ca sunt implementate in kernel si seamana cu procesele, threadurile kernel se mai cheama si *procesuse usoare* (*lightweight processes, LWP*)




Figure 3: Example of kernel thread implementation.

Costuri threaduri kernel

- costul crearii unui thread kernel << costul crearii unui proces
 - se alocă și initializează doar TCB-ul și stivele
 - restul contextului există deja creat în PCB
- context switch-ul între două threaduri **ale aceluiași proces** durează << schimbarea contextului a două procese, pt că nu trebuie schimbat nimic din PCB
- context switch-ul între două threaduri din procese **diferite** are același cost ca și context switch-ul de procese (nu se schimbă doar TCB-urile, ci și PCB-urile)

Planificarea kernel threadurilor pt executie

- threadurile ruleaza asincron unele fata de celelalte si pot pierde procesorul la fel ca si procesele
=> accesul la datele partajate trebuie sincronizat cand se doreste IPC
- planificatorul kernel alege urmatorul thread care trebuie sa ruleze => daca aplicatia are propria politica de planificare (de ex bazata pe prioritati) trebuie sa o comunice intr-un fel sau altul kernelului
- in cazul multiprocesoarelor, planificatorul poate asigna mai multe CPU-uri unui singur proces pt. ca threadurile sa ruleze in paralel ("gang scheduling")
- daca un thread se blocheaza in kernel (de ex prin apel de sistem I/O) si cuanta de timp alocata procesului nu a expirat, planificatorul cauta in lista de TCB-uri un thread gata de rulare *din acelasi proces* si ii acorda procesorul

Protectia kernel threadurilor

- observatia generala despre threaduri e valabila si pt threaduri kernel
- un thread poate corupe stiva altui thread, de pilda => se distrug datele private ale altui thread (variabilele automate/locale)
- solutie: implementarea stivelor in spatii de adresa diferite, dar asta mareste costul context switch-ului
- mai exact, ar fi nevoie de salvarea si reincarcarea contextelor de executie referitoare la gestiunea memoriei, pe langa contextul uzual din TCB

Dezavantajele kernel threadurilor

- desi mai putin costisitoare ca procesele, anumite aspecte le fac nepotrivite pt utilizatori
- (1) *thread_create* este apel sistem => costisitor pt procese care creeaza multe threaduri
- (2) context switching-ul de threaduri necesita intrarea si iesirea in/din kernel mode => overhead aditional context switching-ului obisnuit
- (3) implementarea in kernel e **inflexibila**
 - impune un model de threaduri care nu e potrivit pt orice aplicatie
 - codul planificatorului (scheduler) nu e accesibil (fiind in kernel) => greu de adaptat pt cerintele specifice ale unei anumite aplicatii (politica de planificare nu se poate schimba usor)

User-level threads

- daca planificatorul si TCB-urile sunt implementate in spatiul utilizator costurile scad pentru ca nu mai e nevoie de transgresarea granitei kernel/user
- comparatie calitativa, intr-un ex in care un apel de procedura cost 7 usec, iar un apel sistem (trap) 19 usec:

Operatie	Thread user	Thread kernel	Proces Unix
fork	34 usec	948 usec	11300 usec
IPC synch	37 usec	441 usec	1840 usec

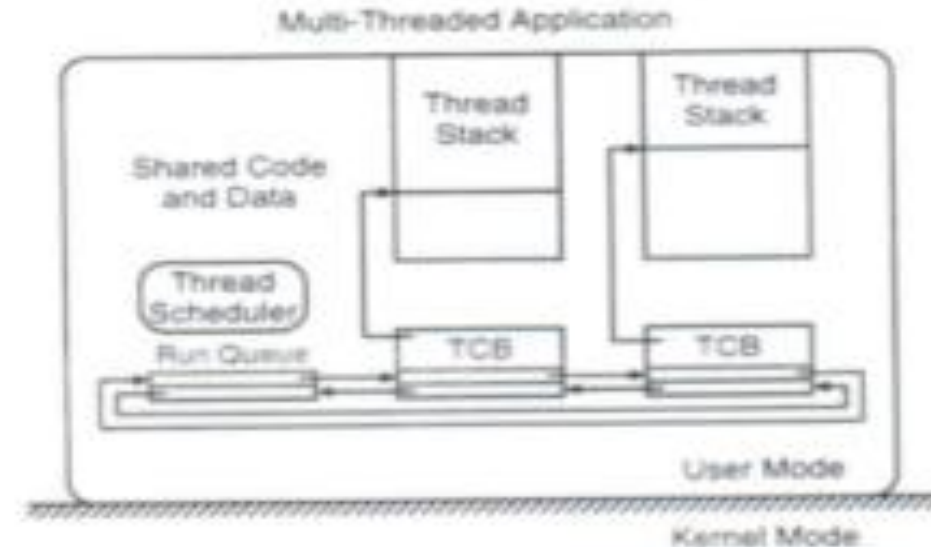


Figure 4: Example of user-level thread implementation.

Caracteristici threaduri utilizator


- nu apeleaza serviciile kernel pt creare si context switch => aceste operatii sunt f. rapide (nu se trece granita user-kernel si nu e nevoie de verificare parametrilor apelului sistem de catre kernel)
- aplicatiile pot furniza propriul planificator (*thread scheduler*) customizat cf unor cerinte specifice
- nu necesita nici un fel de suport explicit din partea kernelului; procesul e privit ca un procesor virtual pt. threaduri
- fiind mult mai rapide decat threadurile kernel, de regula threadurile user se implementeaza deasupra threadurilor kernel


=> planificatorul de threaduri user trateaza threadurile kernel ca pe procesoare virtuale si multiplexeaza mai multe threaduri user pe unul sau mai multe threaduri kernel



Multithreading Models


- Many-to-One
- One-to-One
- Many-to-Many





Many-to-One


- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
 - **Solaris Green Threads**
 - **GNU Portable Threads**





One-to-One


- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
 - Windows
 - Linux






Many-to-Many Model


- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Windows with the *ThreadFiber* package
- Otherwise not very common





Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread




Probleme comune threadurilor kernel si user

- reentranta: multe apeluri de biblioteca sunt ne-reentrante (apelurile reentrante sunt desemnate in paginile de manual ca fiind “thread safe”)
 - ex: trimiterea unui mesaj in retea in 2 pasi: (a) asamblare mesaj in buffer + (b) transmisie (syscall)
 - pierderea procesorului intre (a) si (b) poate conduce la suprascrierea bufferului
 - alte ex: errno, malloc, apeluri stdio
- tratarea semnalelor
 - ex: un thread trateaza un semnal in vreme ce alt thread vrea ca semnalul respectiv sa termine aplicatia (procesul, mai exact)
 - se poate intampla daca se folosesc apeluri de biblioteca si runtime user impreuna
 - problema deriva din faptul ca semnalele sunt definite per proces si nu per thread


Dezavantaje threaduri utilizator

- determinate de faptul ca existenta lor e necunoscuta kernelului
- (1) daca un thread user executa un apel sistem care se blocheaza in kernel, planificatorul kernel (agnostic cu privire la threadurile user) considera ca intreg procesul s-a blocat si alocata CPU altui proces (sau kernel thread), *chiar daca procesul curent mai are si alte threaduri user gata de rulare si nu si-a consumat toata cuanta de timp CPU*
- (2) daca un thread user comite page fault (acces la pagina de memorie inexistentă/nealocata) => acelasi efect ca mai sus, planificatorul kernel alege alt processor/kernel thread in vreme ce pagina de memorie e adusa de pe disc => aplicatia ruleaza cu mai putine CPU decat e necesar
- (3) nefiind constient de existenta threadurilor user, kernelul poate lua procesorul unui thread user care detine un spinlock pe care nu l-a eliberat => scadere dramatica de performanta pt aplicatiile care folosesc threaduri user paralele
 - efectul e si mai dramatic daca schedulerul kernel alege sa ruleze alte threaduri care vor sa obtina acelasi spinlock
- problema de fond: lipsa de coordonare intre schedulerul kernel si implementarea pachetului de threaduri user



Scheduler Activations

- Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**
 - Appears to be a virtual processor on which process can schedule user thread to run
 - Each LWP attached to kernel thread
 - How many LWPs to create?
- Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the **upcall handler** in the thread library
- This communication allows an application to maintain the correct number kernel threads



Scheduler activations

- metoda de coordonare kernel – pachet de threaduri utilizator
- model: aplicatia ruleaza pe un multiprocesor virtual
 - exista apeluri sistem pt a suplimenta/diminua nr de procesoare virtuale alocate de kernel (“adauga CPU”, “acest CPU este idle”)
 - kernelul decide daca onoreaza cererea sau nu
- scheduler activation e aproximarea unui kernel thread
 - are stive kernel si user
 - ofera context de executie pt un thread user
 - in plus, ofera conceptul de *upcall* pt evenimente de mai multe tipuri
 - fiecare upcall creeaza o noua activare (optimizare: folosirea vechilor activari)

Tipuri de evenimente upcall

- adauga procesor – consecinta este executia unui thread user
- procesor preemptat – adauga threadul user care se executa in activarea care a pierdut CPU in coada de threaduri gata de rulare
- activare blocata – activarea s-a blocat (eg, apel sistem blocant) si nu mai utilizeaza procesorul
- activare deblocata – pune in lista gata de rulare threadul care se executa in contextul activarii blocate

Cum functioneaza scheduler activations

- la pornirea procesului
 - kernelul aloca o activare + notifica aplicatia (upcall “adauga CPU”) dupa ce i-a asigurat un CPU
 - sistemul de gestiune al threadurilor user primeste notificarea si foloseste noua activare drept context de executie pt initializarea sa si a threadului *main* (ulterior se pot crea noi threaduri si cere noi procesoare)
- la cerere de suplimentare a concurentei (la adaugarea CPU sau pt. ca un thread se blocheaza)
 - kernelul salveaza starea threadului user in activarea curenta
 - aloca o noua activare si cheama aplicatia in contextul noii activari
- la blocarea unui thread user in kernel
 - nu se continua in kernel la deblocare
 - se creeaza o noua activare
 - se notifica aplicatia
 - schedulerul user copiaza starea threadului blocat din vechea activare si notifica kernelul ca vechea activare poate fi refolosita

Cum functioneaza scheduler activations

- cand un thread user care detine un spinlock pierde procesorul
 - kernelul genereaza un upcall
 - sistemul de gestiune a threadurilor verifica daca threadul e in sectiune critica
 - in caz afirmativ, threadul este continuat temporar printr-un context switch in user space
 - la iesirea din sectiunea critica, se da controlul inapoi upcall-ului tot prin context switch in user space

Exemple de pachete de threaduri user

- Solaris si POSIX
- threadurile Solaris
 - exista si la nivel de kernel (s.n. LWPs, sunt multiplexate pe CPU existente) si la nivel de utilizator
 - pachetul de threaduri (*libthread*) planifica threadurile user pe o colectie de LWPs
 - uzual, threadurile user sunt create nelegate (*unbound*), adica se pot muta intre LWP-uri si asa se si folosesc in general
 - daca un thread user e asignat unui LWP se spune ca e legat (*bound*)
 - controlul kernelului asupra threadurilor user se exercita prin bound threads cu ajutorul setarii prioritatii LWP-ului asociat (de ex, thread user pt stream audio cu prioritate mare in aplicatie multimedia)
 - biblioteca de threaduri asigura ca exista suficiente LWP-uri active pt ca procesul sa poata continua
 - daca un proces determina toate LWP-urile sale sa se blocheze, biblioteca va crea LWP-uri aditionale pt. ca threadurile neblocaate sa poata rula
 - daca LWP-urile sunt idle prea mult timp, biblioteca le da inapoi kernelului pt a fi folosite de alte aplicatii

Pachete de threaduri (cont'd)

- exista apel de setare explicita a nr de LWP-uri la un nou nivel
 - `thr_setconcurrency(int new_level)`
- mecanisme IPC: mutex si variabile conditie
- implementeaza TLS (Thread-Local Storage), capacitate privata de stocare a variabilelor in afara stivei
 - variabile locale nepartajate, “statice” (globale pt threadul respectiv)
 - declarate cu `#pragma`, de ex `#pragma unshared errno;`
- threaduri POSIX (pthreads)
 - standard IEEE, Portable Operating System Interface (e o specificatie, nu o implementare)
 - asemanatoare threadurilor Solaris

Exemplu de pachete de threaduri user

```
int thr_create(void *stack_base, size_t stack_size,
               void *(*function)(void *), void *arg, long flags)
void thr_exit(void *status)
int thr_join(thread_t wait_for, thread_t *joiner_id, void *status)
int mutex_init(mutex_t *mutex, int type, void *arg)
int mutex_lock(mutex_t *mutex)
int mutex_unlock(mutex_t *mutex)
int cond_init(cond_t *cond, int type, int arg)
int cond_wait(cond_t *cond, mutex_t *mutex)
int cond_signal(cond_t *cond)
int cond_broadcast(cond_t *cond)
```

Figure 5: Some commonly used UI thread operations.

```
int pthread_create(pthread_t *thread_pointer,
                   pthread_attr_t *attributes,
                   void *(*function)(void *), void *arg)
void pthread_exit(void *status)
int pthread_join(pthread_t thread, void **status)
void pthread_mutex_init(pthread_mutex_t *mutex,
                        pthread_mutexattr_t *attr)
int pthread_mutex_lock(pthread_mutex_t *mutex)
int pthread_mutex_unlock(pthread_mutex_t *mutex)
int pthread_cond_init(pthread_cond_t *cond, pthread_condattr_t *attr)
int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex)
int pthread_cond_signal(pthread_cond_t *cond)
int pthread_cond_broadcast(pthread_cond_t *cond)
```

Figure 6: Some commonly used Pthreads operations.

```

#include <stdio.h>
#include <errno.h>
#include <pthread.h>

#define ITEMS      8
#define BUFFER_SIZE 4
int buffer[BUFFER_SIZE];

void *producer(void *);
void *consumer(void *);

int spaces;
int items;
int tail;
pthread_cond_t space;
pthread_cond_t item;
pthread_mutex_t buffer_mutex;

char *cmd;

int
main(int argc, char *argv[])
{
    pthread_t producer_thread;
    pthread_t consumer_thread;
    void *thread_return;
    int result;

    cmd = argv[0];

    spaces = BUFFER_SIZE;
    items = 0;
    pthread_mutex_init(&buffer_mutex, NULL);
    pthread_cond_init(&space, NULL);
    pthread_cond_init(&item, NULL);

    if (pthread_create(&producer_thread, NULL, producer, NULL) ||
        pthread_create(&consumer_thread, NULL, consumer, NULL)) {
        fprintf(stderr, "%s: can't make thread\n", cmd);
        exit(1);
    }

    if (pthread_join(producer_thread, &thread_return)) {
        fprintf(stderr, "%s: can't join producer\n", cmd);
        exit(1);
    } else {
        printf("producer returns with %d\n", (int)thread_return);
    }

    if (pthread_join(consumer_thread, &thread_return)) {
        fprintf(stderr, "%s: can't join consumer\n", cmd);
        exit(1);
    } else {
        printf("consumer returns with %d\n", (int)thread_return);
    }

    exit(0);
}

```

Figure 7: Main body of producer/consumer code.

```

void *
producer(void *arg)
{
    int i;

    for (i = 0; i < ITEMS; i++) {
        pthread_mutex_lock(&buffer_mutex);
        while (spaces == 0) {
            pthread_cond_wait(&space, &buffer_mutex);
        }
        buffer[(tail + items) % BUFFER_SIZE] = i;
        items += 1;
        spaces -= 1;
        fprintf(stderr, "producer: puts %d\n", i);
        pthread_mutex_unlock(&buffer_mutex);
        pthread_cond_signal(&item);
    }
    pthread_exit(0);
    return (void *)0;
}

void *
consumer(void *arg)
{
    int i, b;

    for (i = 0; i < ITEMS; i++) {
        pthread_mutex_lock(&buffer_mutex);
        while (items == 0) {
            pthread_cond_wait(&item, &buffer_mutex);
        }
        b = buffer[tail % BUFFER_SIZE];
        tail += 1;
        items -= 1;
        spaces += 1;
        fprintf(stderr, "consumer: gets %d\n", b);
        pthread_mutex_unlock(&buffer_mutex);
        pthread_cond_signal(&space);
    }
    pthread_exit(0);
    return (void *)0;
}


```

Figure 8: Code for producer and consumer threads.



Fork-Join Parallelism

- Multiple threads (tasks) are **forked**, and then **joined**.






Fork-Join Parallelism

- General algorithm for fork-join strategy:

```
Task(problem)
  if problem is small enough
    solve the problem directly
  else
    subtask1 = fork(new Task(subset of problem))
    subtask2 = fork(new Task(subset of problem))



    result1 = join(subtask1)
    result2 = join(subtask2)


    return combined results
```





Fork-Join Parallelism






Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
 - Separating task to be performed from mechanics of creating task allows different strategies for running task
 - ▶ i.e, Tasks could be scheduled to run periodically
- Windows API supports thread pools:


```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
    * this function runs as a separate thread.  
    */  
}
```





Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Signal handling
 - Synchronous and asynchronous
- Thread cancellation of target thread
 - Asynchronous or deferred
- Thread-local storage
- Scheduler Activations






Semantics of `fork()` and `exec()`


- Does `fork()` duplicate only the calling thread or all threads?
 - Some UNIXes have two versions of `fork`
- `exec()` usually works as normal – replace the running process including all threads





Signal Handling

- **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- A **signal handler** is used to process signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled by one of two signal handlers:
 1. default
 2. user-defined
- Every signal has **default handler** that kernel runs when handling signal
 - **User-defined signal handler** can override default
 - For single-threaded, signal delivered to process





Signal Handling (Cont.)

- Where should a signal be delivered for multi-threaded?
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process



Tratarea semnalelor in pachetele de threaduri

- fiecare thread are propria masca de semnale
- toate threadurile unui proces partajeaza handlerele de semnal
- cand SIG_IGN/SIG_DFL sunt setate, se aplica tuturor threadurilor procesului
- semnalele de tip trap (sincrone cu executia threadului, ex SIGFPE, SIGSEGV) executate exclusiv de threadul care le-a generat => mai multe threaduri pot genera si trata acelasi tip de semnal simultan
- semnalele de tip interupere (generate asincron cu executia threadului, ex SIGIO, SIGINT) pot fi tratate de orice thread care nu mascheaza/blocheaza semnalul respectiv (in cazul mai multor astfel de threaduri, se alege unul dintre ele pt tratarea semnalului)
- daca toate threadurile au mascat un anume semnal, se asteapta dupa primul thread care deblocheaza tratarea semnalului respectiv

Tratarea semnalelor in pachetele de threaduri (cont'd)

- `thread_kill`: trimite un semnal care un thread din acelasi proces
 - semnalul e de tip trap, tratat doar de threadul caruia ii este adresat
 - nu se pot trimite semnale unui *anume* thread dintr-un *alt* proces
- threadurile Solaris implementeaza un nou semnal, SIGWAITING (SIG_IGN implicit)
 - trimis procesului cand toate LWP-urile sunt blocate indefinit in asteptarea unui eveniment extern
 - poate fi folosit pt crearea de noi LWP-uri pt a evita blocarea intregului proces
 - idee similara cu scheduler activations, dar coarse-grain (nu merge pt short-term blocking, doar pt evenimente de tip page faults, filesystem I/O)



Thread Cancellation

- Terminating a thread before it has finished
- Thread to be canceled is **target thread**
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
- Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);

/* wait for the thread to terminate */
pthread_join(tid, NULL);
```






Thread Cancellation (Cont.)

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is deferred
 - Cancellation only occurs when thread reaches **cancellation point**
 - ▶ i.e., `pthread_testcancel()`
 - ▶ Then **cleanup handler** is invoked
- On Linux systems, thread cancellation is handled through signals





Linux Threads

- Linux refers to them as **tasks** rather than **threads**
- Thread creation is done through `clone()` system call
- `clone()` allows a child task to share the address space of the parent task (process)
 - Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- `struct task_struct` points to process data structures (shared or unique)

